

employees.csv

emp_id,emp_name,department,job_title,salary,location,hire_date

101,Ramesh,IT,Data Engineer,95000,Bangalore,2019-06-15

102,Anita,HR,HR Executive,65000,Hyderabad,2020-02-10

103,John,Finance,Finance Manager,110000,New York,2016-09-01

104,Meena,IT,Senior Data Analyst,98000,Bangalore,2018-11-23

105,David,IT,Software Engineer,90000,San Francisco,2021-07-05

106,Priya,HR,HR Business Partner,72000,Hyderabad,2019-03-18

107,Michael,Finance,Account Director,120000,New York,2015-01-12

108,Sneha,IT,Junior Developer,78000,Bangalore,2022-08-09

109,Arjun,IT,Data Architect,115000,Bangalore,2017-05-30

110,Emily,Marketing,Marketing Manager,88000,Chicago,2018-10-17

departments.csv

dept_id,department,manager,budget,location

10,IT,Arjun Rao,500000,Bangalore

20,HR,Anita Sharma,200000,Hyderabad

30,Finance,John Smith,600000,New York

40,Marketing,Emily Wilson,300000,Chicago

Snowflake Cloning – Practical Assignments

Important Rule for Students:

- No CTAS
 - No data reload
 - Only cloning & observation
-

Assignment 1: Base Table Analysis

1.1) Load both datasets into Snowflake tables

```
create table employee (emp_id int primary key, emp_name varchar (50),  
department varchar (50),  
job_title varchar (50),  
salary int,location varchar (20), hire_date date);
```

Results 1 of 1 Row • Cluster —

Export Results

	status
1	Table EMPLOYEE successfully created.

create table departments

```
(dept_id int,department varchar (20),  
manager varchar (50),  
budget int,  
location varchar (50));
```

Results 1 of 1 Row • Cluster —

Export Results

	status
1	Table DEPARTMENTS successfully created.

INSERT INTO employee

```
(emp_id, emp_name, department, job_title, salary, location, hire_date)  
VALUES  
(101, 'Ramesh', 'IT', 'Data Engineer', 95000, 'Bangalore','2019-06-15'),  
(102, 'Anita', 'HR', 'HR Executive', 65000, 'Hyderabad', '2020-02-10'),  
(103, 'John', 'Finance', 'Finance Manager',110000, 'New York', '2016-09-01'),  
(104, 'Meena', 'IT', 'Senior Data Analyst', 98000, 'Bangalore', '2018-11-23'),  
(105, 'David', 'IT', 'Software Engineer', 90000, 'San Francisco', '2021-07-05'),  
(106, 'Priya', 'HR', 'HR Business Partner', 72000, 'Hyderabad', '2019-03-18'),  
(107, 'Michael', 'Finance', 'Account Director', 120000, 'New York', '2015-01-12'),  
(108, 'Sneha', 'IT', 'Junior Developer', 78000, 'Bangalore', '2022-08-09'),  
(109, 'Arjun', 'IT', 'Data Architect',115000, 'Bangalore', '2017-05-30'),  
(110, 'Emily', 'Marketing', 'Marketing Manager', 88000, 'Chicago','2018-10-17');
```

Results 1 of 1 Row • Cluster 1

Export Results

	number of rows inserted
1	10

```

INSERT INTO departments
(dept_id, department, manager, budget, location)
VALUES
(10, 'IT', 'Arjun Rao', 500000, 'Bangalore'),
(20, 'HR', 'Anita Sharma', 200000, 'Hyderabad'),
(30, 'Finance', 'John Smith', 600000, 'New York'),
(40, 'Marketing', 'Emily Wilson', 300000, 'Chicago');

```

Results 1 of 1 Row • Cluster 1		Export Results
		number of rows inserted 4

1.2) Validate row counts and data distribution

```

select 'EMPLOYEE' AS TABLE_NAME, COUNT (*) AS ROW_COUNT FROM EMPLOYEE
UNION ALL
SELECT 'DEPARTMENTS' AS TABLE_NAME, COUNT (*) AS ROW_COUNT FROM DEPARTMENTS;

```

Results 2 of 2 Rows • Cluster 1		Export Results
	TABLE_NAME	ROW_COUNT
1	EMPLOYEE	10
2	DEPARTMENTS	4

```

SELECT
department, COUNT (*) AS emp_count,
AVG (salary) AS avg_salary
FROM employee GROUP BY department ORDER BY emp_count DESC;

```

	DEPARTMENT	EMP_COUNT	AVG_SALARY
1	IT	5	95200.000000
2	HR	2	68500.000000
3	Finance	2	115000.000000
4	Marketing	1	88000.000000

```
select job_title,location,Avg(salary)as avg_salary from employee group by location,job_title,emp_name order by avg_salary Desc;
```

Results		10 of 10 Rows • Cluster 1	Export Results
	JOB_TITLE	LOCATION	AVG_SALARY
1	Account Director	New York	120000.00000
2	Data Architect	Bangalore	115000.00000
3	Finance Manager	New York	110000.00000
4	Senior Data Analyst	Bangalore	98000.00000
5	Data Engineer	Bangalore	95000.00000
6	Software Engineer	San Francisco	90000.00000
7	Marketing Manager	Chicago	88000.00000
8	Junior Developer	Bangalore	78000.00000
9	HR Business Partner	Hyderabad	72000.00000
10	HR Executive	Hyderabad	65000.00000

1.3) Identify candidate tables for cloning

Candidate table => the **original/source table** that we *choose* for cloning

Candidate tables were identified based on data availability, frequency of changes, business relevance, and suitability for observing Snowflake cloning behavior.

```
SELECT table_name, row_count  
FROM TAB_DB.INFORMATION_SCHEMA.TABLES  
WHERE table_schema = 'PUBLIC'  
AND table_name IN ('EMPLOYEE', 'DEPARTMENTS');
```

INFORMATION_SCHEMA is a system schema that provides metadata about tables, schemas, and other database objects in Snowflake.

Res (14 hours ago) | Results (just now)

Table Chart Q ⚡ 2 rows ⓘ 1.8s ⏪ ⏴

#	TABLE_NAME	# ROW_COUNT
1	DEPARTMENTS	4
2	EMPLOYEE	10

1.4) Explain why these tables are suitable for cloning experiments:

These tables are suitable for cloning because Employees supports frequent data changes while Departments represents stable reference data. Their size and business relevance make them ideal for demonstrating Snowflake's zero-copy cloning and isolation behavior.

Assignment 2: Table-Level Cloning

2.1) Clone the employee table

```
create table clone_emp clone employee;
```

Results 1 of 1 Row • Cluster —	
	status
1	Table CLONE_EMP successfully created.

2.2) Compare:

Row count:

```
select count (*) from employee;
```

Results 1 of 1 Row • Cluster —	
	COUNT(*)
1	10

```
select count (*) from clone_emp;
```

Results	
Results (just now)	
Table	Chart
1	## EMPCLONE_ROW_COUNT 10

Query Performance:

```
SELECT department,COUNT(*) AS emp_count,AVG(salary) AS avg_salary
```

```
FROM employee GROUP BY department
```

```
ORDER BY emp_count DESC;
```

Results	
Results (just now)	
Table	Chart
1	## DEPARTMENT
2	## EMP_COUNT
3	## AVG_SALARY
1	IT
2	5
2	HR
2	2
3	Finance
2	115000.000000
4	Marketing
1	88000.000000

```
SELECT department,COUNT(*) AS emp_count,AVG(salary) AS avg_salary
```

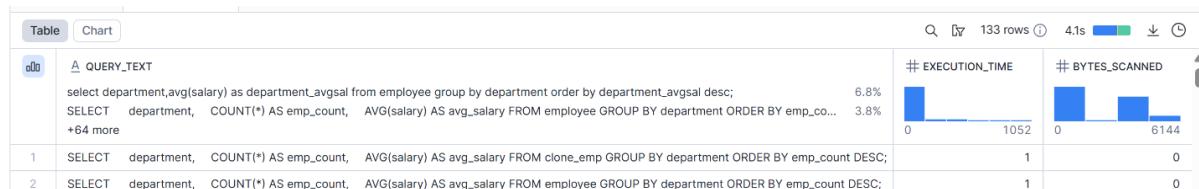
```
FROM clone_emp GROUP BY department
```

```
ORDER BY emp_count DESC;
```

Results (just now)

Table

	DEPARTMENT	# EMP_COUNT	# AVG_SALARY
1	IT	5	95200.000000
2	HR	2	68500.000000
3	Finance	2	115000.000000
4	Marketing	1	88000.000000



Storage usage is checked to verify Snowflake's zero-copy cloning behavior. Identical storage values for the original and cloned tables confirm that no physical data is duplicated during clone creation. Storage increases only when data is modified, demonstrating the copy-on-write mechanism.

2.3) Explain why clone creation is instantaneous

Snowflake clone creation is instantaneous because only metadata is copied, not the actual data. Data is physically duplicated only when modifications occur through copy-on-write.

Assignment 3: Change Isolation Testing

3.1) Update salary data in the original employee table:

UPDATE employee

SET salary = salary + 1000

WHERE department = 'IT';

Results (just now)

Table

	# number of rows updated	# number of multi-joined rows updated
1	5	0

--validate

SELECT * FROM employee where DEPARTMENT='IT';

Results (just now)

Table

	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
1	101	Ramesh	IT	Data Engineer	96000	Bangalore	2019-06-15
2	104	Meena	IT	Senior Data Analyst	99000	Bangalore	2018-11-23
3	105	David	IT	Software Engineer	91000	San Francisco	2021-07-05
4	108	Sneha	IT	Junior Developer	79000	Bangalore	2022-08-09
5	109	Arjun	IT	Data Architect	116000	Bangalore	2017-05-30

Eg: Before 95000 and after 96000

Select * from clone_emp WHERE DEPARTMENT='IT';

EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
101	Ramesh	IT	Data Engineer	95000	Bangalore	2019-06-15
104	Meena	IT	Senior Data Analyst	98000	Bangalore	2018-11-23
105	David	IT	Software Engineer	90000	San Francisco	2021-07-05
108	Sneha	IT	Junior Developer	78000	Bangalore	2022-08-09
109	Arjun	IT	Data Architect	115000	Bangalore	2017-05-30
115	anuj	IT	Data Architect	115000	Bangalore	2017-05-30

Employee table IT department only changed

Its not affected in clone_emp

3.2) Delete records in the cloned table

DELETE FROM CLONE_EMP WHERE EMP_ID = 102;

Results 1 of 1 Row • Cluster 1		Export Results
		number of rows deleted
1		1

Validate

select count(emp_id)from clone_emp where emp_id ='102' order by emp_id;

Results 1 of 1 Row • Cluster 1		Export Results
		COUNT(EMP_ID)
1		0

SELECT count(emp_id) FROM employee where emp_id= '102' order by emp_id;

Results 1 of 1 Row • Cluster 1		Export Results
		COUNT(EMP_ID)
1		1

Once we delete the clone table data only deleted in clone data its not affected the primary table data.

3.3)Insert new records in the clone

insert into clone_emp values(115, 'anuj', 'IT', 'Data Architect', 115000, 'Bangalore', '2017-05-30'),
(116, 'miky', 'Marketing', 'Marketing Manager', 88000, 'Chicago', '2018-10-17');

Results		Table	Chart	Results (just now)			
						1 row	232ms
000	# number of rows inserted						
1						2	

--Validate

select count(emp_id) from employee;

Results		Table	Chart	Results (just now)			
						1 row	74ms
000	# COUNT(EMP_ID)						
1						10	

select count(emp_id) from clone_emp;

Results		Table	Chart	Results (just now)			
						1 row	123ms
000	# COUNT(EMP_ID)						
1						11	

I insert the data in clone table and successfully inserted but employee table is not affected.

3.3) Validate isolation between original and clone

Operation	Original	Clone
Updated in original table	Its Changed	Its not change
Updated in clone table	Its not affected	Its changed
Insert in original	Its inserted in original only	Not inserted in clone table
Insert in clone	its not inserted	Clone table only inserted
Delete in original	Only original table data is deleted	Clone table its not affected
Delete in clone	Its not changed	Its deleted

3.5) Explain backend copy-on-write behavior

Snowflake uses a copy-on-write mechanism for cloned objects. Initially, both the original and cloned tables share the same underlying storage. When a DML operation such as UPDATE, DELETE, or INSERT is performed, Snowflake creates new micro-partitions only for the modified data. Unchanged data continues to be shared. This ensures complete isolation between the original table and the clone while maintaining storage efficiency.

4: Storage Growth Exploration

Track storage:

```
SELECT table_name, row_count, bytes  
FROM TAB_DB.INFORMATION_SCHEMA.TABLES  
WHERE table_schema = 'PUBLIC'  
AND table_name = 'EMPLOYEE';
```

Results			
Results (just now)			
Table		Chart	
#	TABLE_NAME	# ROW_COUNT	# BYTES
1	EMPLOYEE	10	3072

4.1) Perform DML operations incrementally:

```
insert into employee values (118, 'Emily', 'Marketing', 'Marketing Manager', 88000, 'Chicago', '2018-10-17');
```

```
insert into employee values (119, 'Eoy', 'Marketing', 'Marketing Manager', 88000, 'Chicago', '2018-10-17');
```

Results			
Results (just now)			
Table		Chart	
#	# number of rows inserted		
1			1

Results			
Results (just now)			
Table		Chart	
#	# number of rows inserted		
1			1

Performance Check:

```
SELECT table_name, row_count, bytes  
FROM TAB_DB.INFORMATION_SCHEMA.TABLES  
WHERE table_schema = 'PUBLIC'  
AND table_name = 'EMPLOYEE';
```

Results			
Results (just now)			
Table		Chart	
#	TABLE_NAME	# ROW_COUNT	# BYTES
1	EMPLOYEE	10	3072

4.2) Track storage before and after each operation

Update the table:

```
update employee set salary = salary + 10000 where emp_id = 105;
```

Results		Results (just now)	
		Table Chart	
#	number of rows updated	#	number of multi-joined rows updated
1		1	0

Track storage:

```
SELECT table_name, row_count, bytes
FROM TAB_DB.INFORMATION_SCHEMA.TABLES
WHERE table_schema = 'PUBLIC'
AND table_name = 'EMPLOYEE';
```

Results		Results (1 minute ago)	
		Table Chart	
#	TABLE_NAME	# ROW_COUNT	# BYTES
1	EMPLOYEE	10	3072

Delete the record in employee table:

```
delete from employee where emp_id = 102;
```

Results		Results (just now)	
		Table Chart	
#	number of rows deleted		
1			1

```
delete from employee where emp_id=103;
```

Results		1 of 1 Row • Cluster 1	
		Export Results	
1		number of rows deleted	1

Track Storage:

```
SELECT table_name, row_count, bytes
FROM TAB_DB.INFORMATION_SCHEMA.TABLES
WHERE table_schema = 'PUBLIC'
AND table_name = 'EMPLOYEE';
```

Results		Results (just now)	
		Table Chart	
#	TABLE_NAME	# ROW_COUNT	# BYTES
1	EMPLOYEE	10	3072

Update the 5 Records:

```
UPDATE employee
```

```
SET salary = salary + 5000
```

```
WHERE emp_id IN (105, 106, 108, 109, 110);
```

Results 1 of 1 Row • Cluster 1		Export Results
	number of rows updated	number of multi-joined rows updated
1	5	0

Storage Track Progress:

```
SELECT table_name, row_count,  
Bytes FROM TAB_DB.INFORMATION_SCHEMA.TABLES  
WHERE table_schema = 'PUBLIC'  
AND table_name = 'EMPLOYEE';
```

Results 1 of 1 Row • Cluster 1		Export Results
TABLE_NAME	ROW_COUNT	BYTES
EMPLOYEE	10	3072

Update Employee Table:

```
UPDATE employee
```

```
SET salary = salary + 1000;
```

Results 1 of 1 Row • Cluster 1		Export Results
	number of rows updated	number of multi-joined rows updated
1	10	0

Results 1 of 1 Row • Cluster 1		Export Results
TABLE_NAME	ROW_COUNT	BYTES
EMPLOYEE	10	3072

During the experiment, storage usage did not visibly increase because the table contained a very small dataset. Snowflake stores data in immutable micro-partitions, and small DML changes were absorbed within existing partitions without creating new ones.

4.3) Identify the exact moment storage increases

In Snowflake, storage growth occurs during UPDATE operations because the copy-on-write mechanism creates new micro-partitions instead of modifying existing data.

4.4) Explain Snowflake's physical storage strategy

Snowflake stores data in compressed, columnar **micro-partitions** with rich metadata.

Data is **immutable**, so changes create new micro-partitions instead of modifying existing ones.

This design enables **copy-on-write, zero-copy cloning, and Time Travel** with efficient storage use.

Assignment 5: Time Travel Cloning

5.1) Perform multiple updates on the original table

`SELECT * FROM employee;`

Results		Results (just now)									
#	EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE				
	101	Emily Eoy +8 more	IT Marketing +2 more	Marketing Manager Account Director +7 more	78000 121000	Bangalore Chicago +3 more	1/12/2015 8/9/2022				
1	101	Ramesh	IT	Data Engineer	96000	Bangalore	2019-06-15				
2	103	John	Finance	Finance Manager	111000	New York	2016-09-01				
3	104	Meena	IT	Senior Data Analyst	99000	Bangalore	2018-11-23				
4	107	Michael	Finance	Account Director	121000	New York	2015-01-12				
5	105	David	IT	Software Engineer	106000	San Francisco	2021-07-05				
6	106	Priya	HR	HR Business Partner	78000	Hyderabad	2019-03-18				
7	108	Sneha	IT	Junior Developer	84000	Bangalore	2022-08-09				

`update employee set salary=salary+5000 where emp_id=101;`

`update employee set salary = salary +5000 where emp_id=104;`

`update employee set salary = salary +5000 where emp_id=119;`

`update employee set salary = salary +5000 where emp_id=110;`

`update employee set salary = salary +5000;`

Results		Results (1 minute ago)									
#	EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE				
	101	Emily Eoy +8 more	IT Marketing +2 more	Marketing Manager Account Director +7 more	83000 126000	Bangalore Chicago +3 more	1/12/2015 8/9/2022				
1	101	Ramesh	IT	Data Engineer	106000	Bangalore	2019-06-15				
2	103	John	Finance	Finance Manager	116000	New York	2016-09-01				
3	104	Meena	IT	Senior Data Analyst	109000	Bangalore	2018-11-23				
4	107	Michael	Finance	Account Director	126000	New York	2015-01-12				
5	105	David	IT	Software Engineer	111000	San Francisco	2021-07-05				
6	106	Priya	HR	HR Business Partner	83000	Hyderabad	2019-03-18				
7	108	Sneha	IT	Junior Developer	89000	Bangalore	2022-08-09				

After Updated the multiple records

5.2) Create clones from:

Current state:

`CREATE OR REPLACE TABLE employee_clone_curr --we update current records that clone table`

`CLONE employee;`

Results		Results (just now)		
			Q	1 row ⓘ 1.3s ⚡ ⬇ ⓘ
Table		Chart		
EMPLOYEE_CLONE_CURR	A status			
1	Table EMPLOYEE_CLONE_CURR successfully created.			

select*from employee_clone_curr;

Results		Results (just now)		
			Q	12 rows ⓘ 428ms ⚡ ⬇ ⓘ
Table		Chart		
EMPLOYEE_CLONE_CURR	# EMP_ID	A EMP_NAME	A DEPARTMENT	A JOB_TITLE
100	101	Emily	IT	Marketing Manager
	102	Eoy	Marketing	Account Director
	103	+8 more	+3 more	+7 more
1	101	Ramesh	IT	Data Engineer
2	103	John	Finance	Finance Manager
3	104	Meena	IT	Senior Data Analyst
4	107	Michael	Finance	Account Director
5	105	David	IT	Software Engineer
6	106	Priya	HR	HR Business Partner
7	108	Sneha	IT	Junior Developer
				# SALARY
				83000 126000
				106000 116000 126000
				Bangalore Chicago +3 more
				33.3% 33.3%
				1/12/2015 8/9/2022
				2019-06-15 2016-09-01 2018-11-23 2015-01-12 2021-07-05 2019-03-18 2022-08-09
				HIRE_DATE

Creating a clone from the current state means making an instant copy that reflects the latest data at the time of cloning.

It uses zero-copy cloning, so no data is physically duplicated and future changes stay isolated.

Past state (Time Travel)

SELECT CURRENT_TIMESTAMP();

Results		Results (just now)		
			Q	1 row ⓘ 53ms ⚡ ⬇ ⓘ
Table		Chart		
EMPLOYEE_CLONE_PAST	CURRENT_TIMESTAMP()			
1	2025-12-22 05:45:02.716 -0800			

CREATE OR REPLACE TABLE employee_clone_past

CLONE employee

AT (TIMESTAMP => '2025-12-22 05:00:00'::TIMESTAMP_TZ);

Results		Results (6 minutes ago)		
			Q	1 row ⓘ 952ms ⚡ ⬇ ⓘ
Table		Chart		
EMPLOYEE_CLONE_PAST	A status			
1	Table EMPLOYEE_CLONE_PAST successfully created.			

Select * from employee_clone_past;

Results | Results (just now) | 10 rows 684ms

Table | Chart

	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
	101 110	Anita Arjun +8 more	10.0% 10.0% +2 more	IT Finance	50.0% 20.0%	Account Director Data Architect +8 more	10.0% 10.0%
1	101	Ramesh	IT	Data Engineer	95000	Bangalore	40.0%
2	102	Anita	HR	HR Executive	65000	Hyderabad	20.0%
3	103	John	Finance	Finance Manager	110000	New York	2016-09-01
4	104	Meena	IT	Senior Data Analyst	98000	Bangalore	2018-11-23
5	105	David	IT	Software Engineer	90000	San Francisco	2021-07-05
6	106	Priya	HR	HR Business Partner	72000	Hyderabad	2019-03-18

5.3) Compare data across versions

Original Table

SELECT * FROM employee ORDER BY emp_id;

Results | Results (just now) | 12 rows 30ms

Table | Chart

	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
	101 119	Emily Eoy +8 more	16.7% 16.7% +2 more	IT Marketing	41.7% 33.3%	Marketing Manager Account Director +7 more	33.3% 8.3%
1	101	Ramesh	IT	Data Engineer	106000	Bangalore	2019-06-15
2	103	John	Finance	Finance Manager	116000	New York	2016-09-01
3	104	Meena	IT	Senior Data Analyst	109000	Bangalore	2018-11-23
4	105	David	IT	Software Engineer	111000	San Francisco	2021-07-05
5	106	Priya	HR	HR Business Partner	83000	Hyderabad	2019-03-18
6	107	Michael	Finance	Account Director	126000	New York	2015-01-12
7	108	Sneha	IT	Junior Developer	89000	Bangalore	2022-08-09

SELECT * FROM employee_clone_curr ORDER BY emp_id;

Updated data its only available

Results | Results (just now) | 12 rows 27ms

Table | Chart

	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
	101 119	Emily Eoy +8 more	16.7% 16.7% +2 more	IT Marketing	41.7% 33.3%	Marketing Manager Account Director +7 more	33.3% 8.3%
1	101	Ramesh	IT	Data Engineer	106000	Bangalore	2019-06-15
2	103	John	Finance	Finance Manager	116000	New York	2016-09-01
3	104	Meena	IT	Senior Data Analyst	109000	Bangalore	2018-11-23
4	105	David	IT	Software Engineer	111000	San Francisco	2021-07-05
5	106	Priya	HR	HR Business Partner	83000	Hyderabad	2019-03-18
6	107	Michael	Finance	Account Director	126000	New York	2015-01-12
7	108	Sneha	IT	Junior Developer	89000	Bangalore	2022-08-09

SELECT * FROM employee_clone_past ORDER BY emp_id;

Before updated data its only using Timestamp

Results | Results (just now) | X

Table **Chart**

	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
101	110	Anita Arjun +8 more	IT Finance +2 more	Account Director Data Architect +8 more	65000 120000	Bangalore Hyderabad +3 more	40.0% 20.0% +3 more
1	101	Ramesh	IT	Data Engineer	95000	Bangalore	2019-06-15
2	102	Anita	HR	HR Executive	65000	Hyderabad	2020-02-10
3	103	John	Finance	Finance Manager	110000	New York	2016-09-01
4	104	Meena	IT	Senior Data Analyst	98000	Bangalore	2018-11-23
5	105	David	IT	Software Engineer	90000	San Francisco	2021-07-05
6	106	Priya	HR	HR Business Partner	72000	Hyderabad	2019-03-18
7	107	Michael	Finance	Account Director	120000	New York	2015-01-12

5.4) Explain real-world recovery use cases

Real-time Use Case 1: Accidental Update

Production-la mistake-aa salary update aagiduchu na:

Time Travel clone eduthu

Correct data recover pannalaam

Use Case 2: Data Validation

New transformation apply panna munna:

Old data clone eduthu

Compare pannalaam

Assignment 6: Schema-Level Cloning

6.1) Clone a schema containing employee and department tables

show tables in schema tab_db.public;

Results | Results (just now) | X

Table **Chart**

	created_on	name	database_name	schema_name	kind	comment	cluster_by	rows	bytes	owner
1	12/18/2025	CLONE_EMP	TAB_DB	PUBLIC	TABLE	... 92.3% ... 7.7%	100.0%	4	12	ACCOUNTADMIN 100.
2	2025-12-21 04:13:26.150 -0800	CLONE_TAB	TAB_DB	PUBLIC	TABLE			11	3072	ACCOUNTADMIN
3	2025-12-19 03:12:45.429 -0800	CLONE_TB	TAB_DB	PUBLIC	TABLE			6	2048	ACCOUNTADMIN
4	2025-12-18 21:50:36.239 -0800	CLP	TAB_DB	PUBLIC	TABLE			12	2560	ACCOUNTADMIN
5	2025-12-21 03:26:06.105 -0800	DEPARTMENTS	TAB_DB	PUBLIC	TABLE			10	3072	ACCOUNTADMIN
6	2025-12-21 03:24:16.991 -0800	EMPLOYEE	TAB_DB	PUBLIC	TABLE			4	2560	ACCOUNTADMIN
								10	3072	ACCOUNTADMIN

create or replace schema tab_db.clone_pub_schema clone tab_db.public;

Results | Results (just now) | X

Table **Chart**

	status
1	Schema CLONE_PUB_SCHEMA successfully created.

SHOW TABLES IN SCHEMA TAB_DB.clone_pub_schema;

Table | Chart

13 rows 49ms ACCO

	LT created_on	A name	A database_name	A schema_name	A kind	A comment	A cluster_by	# rows	# bytes	A owner
	12/22/2025 12/22/2025	CLON... 7.7% CLON... 7.7% +11 more	TAB_DB 100.0%	CLONE_PUB_S... 100.0%	... 92.3% ... 7.7%	100.0%	100.0%	4 12	15...30...	ACCO
1	2025-12-22 06:07:39.930 -0800	CLONE_EMP	TAB_DB	CLONE_PUB_SCHEMA	TABLE			11	3072	ACCO
2	2025-12-22 06:07:39.930 -0800	CLONE_TAB	TAB_DB	CLONE_PUB_SCHEMA	TABLE			6	2048	ACCO
3	2025-12-22 06:07:39.930 -0800	CLONE_TB	TAB_DB	CLONE_PUB_SCHEMA	TABLE			12	2560	ACCO
4	2025-12-22 06:07:39.930 -0800	CLP	TAB_DB	CLONE_PUB_SCHEMA	TABLE			10	3072	ACCO
5	2025-12-22 06:07:39.930 -0800	DEPARTMENTS	TAB_DB	CLONE_PUB_SCHEMA	TABLE			4	2560	ACCO
6	2025-12-22 06:07:39.930 -0800	EMPLOYEE	TAB_DB	CLONE_PUB_SCHEMA	TABLE			10	3072	ACCO

6.2)Modify tables in cloned schema

USE SCHEMA TAB_DB.clone_PUB_schema;

UPDATE departments

SET budget = budget + 300000

WHERE dept_id = 20;

Results | Results (just now)

Table | Chart

1 row 1.7s ACCO

	# number of rows updated	# number of multi-joined rows updated
1	1	0

INSERT INTO departments VALUES (201, 'IT', 'Arjun Rao', 500000, 'Bangalore');

Results | Results (just now)

Table | Chart

1 row 1.1s ACCO

	# number of rows inserted
1	1

select*from departments;

Results | Results (just now)

Table | Chart

5 rows 115ms ACCO

	# DEPT_ID	A DEPARTMENT	A MANAGER	# BUDGET	A LOCATION
1	10	IT	Arjun Rao	500000	Bangalore
2	20	HR	Anita Sharma	500000	Hyderabad
3	30	Finance	John Smith	600000	New York
4	40	Marketing	Emily Wilson	300000	Chicago
5	201	IT	Arjun Rao	500000	Bangalore

6.3)Drop tables in original schema

DROP TABLE TAB_DB.PUBLIC.DEPARTMENTS;

Results | Results (just now)

Table | Chart

1 row 89ms ACCO

	A status
1	DEPARTMENTS successfully dropped.

SHOW TABLES IN SCHEMA TAB_DB.PUBLIC;

Its not show

Results | Results (just now) | 12 rows 72ms

Table | Chart

#	created_on	A name	A database_name	A schema_name	A kind	A comment	A cluster_by	# rows	# bytes	A owner
	12/18/2025	CLON...	8.3%			... 91.7%		6	12	
	12/21/2025	CLON...	8.3%	+10 more	TAB_DB	100.0%	PUBLIC	100.0%	15...30...	ACCOUNT... 100.
1	2025-12-21 04:13:26.150-0800	CLONE_EMP	TAB_DB	PUBLIC	TABLE			11	3072	ACCOUNTADMIN
2	2025-12-19 03:12:45.429-0800	CLONE_TAB	TAB_DB	PUBLIC	TABLE			6	2048	ACCOUNTADMIN
3	2025-12-18 21:50:36.239-0800	CLONE_TB	TAB_DB	PUBLIC	TABLE			12	2560	ACCOUNTADMIN
4	2025-12-21 04:35:40.433-0800	CLP	TAB_DB	PUBLIC	TABLE			10	3072	ACCOUNTADMIN
5	2025-12-21 03:24:16.991-0800	EMPLOYEE	TAB_DB	PUBLIC	TABLE			10	3072	ACCOUNTADMIN
6	2025-12-21 07:16:29.482-0800	EMPLOYEE_CLONE	TAB_DB	PUBLIC	TABLE			10	3072	ACCOUNTADMIN
7	2025-12-21 07:23:34.875-0800	EMPLOYEE_CLONE	TAB_DB	PUBLIC	TABLE			10	3072	ACCOUNTADMIN
8	2025-12-19 03:11:15.005-0800	EMP_PERM	TAB_DB	PUBLIC	TABLE			6	2048	ACCOUNTADMIN

```
SELECT * FROM TAB_DB.clone_pub_schema.departments;
```

Results | Results (just now) | 5 rows 69ms

Table | Chart

#	# DEPT_ID	A DEPARTMENT	A MANAGER	# BUDGET	A LOCATION
1	10	IT	Arjun Rao	500000	Bangalore
2	20	HR	Anita Sharma	500000	Hyderabad
3	30	Finance	John Smith	600000	New York
4	40	Marketing	Emily Wilson	300000	Chicago
5	201	IT	Arjun Rao	500000	Bangalore

```
SELECT * FROM TAB_DB.public.departments;
```

Results | ... | Results (just now) | 0 rows 23ms



SQL compilation error: Object 'TAB_DB.PUBLIC.DEPARTMENTS' does not exist or not authorized.

6.4) Observe behavior and explain metadata inheritance

Schema-level cloning in Snowflake creates a logical copy of all tables using metadata inheritance. The cloned schema initially shares the same underlying data. When data is modified in the cloned schema, Snowflake uses copy-on-write to create new micro-partitions. Dropping tables in the original schema does not impact the cloned schema because each clone maintains independent metadata references.

Assignment 7: Database-Level Cloning

7.1) Clone the entire database

show schemas in database tab_db;

id	created_on	name	is_default	is_current	database_name	owner	comment	options	retention_in
1	2025-12-22 06:07:39.930 -0800	CLONE_PUB_SCHEMA	N	Y	TAB_DB	ACCOUNTADMIN			1
2	2025-12-22 08:15:13.014 -0800	INFORMATION_SCHEMA	N	N	TAB_DB		Views describing the c		1
3	2025-12-17 20:49:10.125 -0800	PUBLIC	N	N	TAB_DB	ACCOUNTADMIN			1

show databases;

id	created_on	name	is_default	is_current	origin	owner	comment	options	retent
1	2025-12-10 08:16:13.607 -0800	CLI_DB	N	N		ACCOUNTADMIN			1
2	2025-12-15 08:44:22.856 -0800	SECURITY_DB	N	N		ACCOUNTADMIN			1
3	2025-12-02 06:09:37.790 -0800	SNOWFLAKE	N	N	SNOWFLAKE.ACCTOUN				0
4	2025-12-02 06:09:45.924 -0800	SNOWFLAKE_LEARNING_DB	N	N		ACCOUNTADMIN	Created by Snowflake		1
5	2025-12-02 06:09:43.728 -0800	SNOWFLAKE_SAMPLE_DATA	N	N	SFSALESSHARED.SFC	ACCOUNTADMIN	Preloaded TPCH Data		0
6	2025-12-17 20:49:10.104 -0800	TAB_DB	N	Y		ACCOUNTADMIN			1
7	2025-12-09 00:26:21.721 -0800	TRAINING	N	N		ACCOUNTADMIN			1
8	2025-12-03 21:31:05.725 -0800	USER\$HARIHARAN12	N	N					1

create or replace database clone_tab_db clone tab_db;

status
Database CLONE_TAB_DB successfully created.

7.2) Test query access and object availability

USE DATABASE clone_tab_db;

SELECT * FROM PUBLIC.EMPLOYEE;

Results ... Results (7 minutes ago)

Table Chart

Q 10 rows 150ms ↴ ⏺

#	EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
1	107	Michael	Finance	Account Director	126000	New York	2015-01-12
2	101	Ramesh	IT	Data Engineer	107000	Bangalore	2019-06-15
3	104	Meena	IT	Senior Data Analyst	110000	Bangalore	2018-11-23
4	106	Priya	HR	HR Business Partner	83000	Hyderabad	2019-03-18
5	105	David	IT	Software Engineer	112000	San Francisco	2021-07-05
6	108	Sneha	IT	Junior Developer	90000	Bangalore	2022-08-09
7	109	Arjun	IT	Data Architect	127000	Bangalore	2017-05-30

7.3) Perform DML on cloned database

UPDATE clone_tab_db.PUBLIC.EMPLOYEE

SET salary = salary + 3000 WHERE emp_id = 104;

Results ... Results (just now)

Table Chart

Q 1 row 987ms ↴ ⏺

#	number of rows updated	number of multi-joined rows updated
1	1	0

Clone database Its perform DML commands

insert into clone_tab_db.public.employee values(301, 'Ramesh', 'IT', 'Data Engineer', 95000, 'Bangalore', '2019-06-15');

Results ... Results (just now)

Table Chart

Q 1 row 773ms ↴ ⏺

#	number of rows inserted
1	1

SELECT * FROM TAB_DB.PUBLIC.EMPLOYEE WHERE emp_id = 104; -- 110000 primary database table is not updated

Results ... Results (just now)

Table Chart

Q 1 row 71ms ↴ ⏺

#	EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
1	104	Meena	IT	Senior Data Analyst	110000	Bangalore	2018-11-23

SELECT*FROM CLONE_TAB_DB.PUBLIC.EMPLOYEE WHERE EMP_ID=104

Results ... Results (just now)

Table Chart

Q 1 row 1.4s ↴ ⏺

#	EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
1	104	Meena	IT	Senior Data Analyst	113000	Bangalore	2018-11-23

--113000 ITS CHANGED SO CLONE DATABASE ITS CHANGED

7.4) Explain how Snowflake isolates environments (Dev/Test/Prod)

Snowflake isolates Dev, Test, and Prod environments using database-level cloning. Each cloned database inherits metadata and shares the same underlying data initially. When DML operations are performed, Snowflake uses copy-on-write to isolate changes, ensuring that modifications in one environment do not impact others.

Assignment 8: Privileges & Security

8.1) Analyze access privileges before cloning

SHOW GRANTS ON TABLE EMPLOYEE;

Results (just now)									
Table		Chart							
id	privilege	granted_on	name	granted_to	grantee_name	grant_option	granted_by	granted_by_role_type	
1	OWNERSHIP	TABLE	CLONE_TAB_DB.PUBLIC.EMPLOYEE	ROLE	ACCOUNTADMIN	true	ACCOUNTADMIN	ROLE	

8.2) Observe privileges after cloning

SHOW GRANTS ON TABLE CLONE_EMP;

Results (just now)									
Table		Chart							
id	on	privilege	granted_on	name	granted_to	grantee_name	grant_option	granted_by	granted_by_role_type
1	22 21:25:06	OWNERSHIP	TABLE	CLONE_TAB_DB.PUBLIC.CLONE_EMP	ROLE	ACCOUNTADMIN	true	ACCOUNTADMIN	ROLE

8.3) Modify privileges on clone

SHOW GRANTS ON TABLE CLONE_EMP;

We checked the privileges on the EMPLOYEE table and saw which roles had access, like SELECT for the ANALYST role.

After cloning the table, those privileges were not copied automatically.

Only the owner role has full access on the clone, and any other access must be granted manually. Changes made on the clone do not affect the original table's privileges.

Results (2 minutes ago)									
Table		Chart							
id	created_on	privilege	granted_on	name	granted_to	grantee_name	grant_option	granted_by	granted_by_role_type
1	2025-12-22 21:25:06.267 -0800	OWNERSHIP	TABLE	CLONE_TAB_DB.PUBLIC.CLONE_EMP	ROLE	ACCOUNTADMIN	true	ACCOUNTADMIN	ROLE
2	2025-12-22 23:44:06.849 -0800	SELECT	TABLE	CLONE_TAB_DB.PUBLIC.CLONE_EMP	ROLE	DEVELOPER	false	ACCOUNTADMIN	ROLE

8.4) Explain why Snowflake handles security this way

Snowflake handles security very strictly. Even though cloned tables share the same data, access privileges are not shared. This prevents unauthorized users from accessing cloned data. This design follows the least-privilege principle and helps keep production data secure.

Assignment 9: Performance Comparison

9.1) Run identical analytical queries on:

Original tables:

```
use tab_db;
```

```
select department,avg(salary) as department_avgsal from employee group by department order by department_avgsal desc;
```

#	DEPARTMENT	# DEPARTMENT_AVGSAL
1	Finance	126000.000000
2	IT	109200.000000
3	Marketing	99000.000000
4	HR	83000.000000

Cloned tables:

```
use clone_tab_db;
```

```
select department,avg(salary) as department_avgsal from employee group by department order by department_avgsal desc;
```

#	DEPARTMENT	# DEPARTMENT_AVGSAL
1	Finance	115000.000000
2	IT	98500.000000
3	Marketing	88000.000000
4	HR	72000.000000

9.2) Compare execution plans:

EXPLAIN

```
SELECT department, AVG(salary)as department_avgsal
```

```
FROM TAB_DB.PUBLIC.EMPLOYEE
```

```
GROUP BY department order by department_avgsal desc;
```

#	# id	# parentOperators	# operation	# objects	# alias	# expressions	# partitionsTotal	# partitionsAssigned	# bytesAssigned
1	null	null	GlobalStats	null	null	null	1	1	3072
2	1	0	Result	null	null	EMPLOYEE.DEPARTME	null	null	null
3	1	1	[0]	Sort	null	SCALED_ROUND_INT_C	null	null	null
4	1	2	[1]	Aggregate	null	aggExprs: [SUM(EMPL	null	null	null
5	1	3	[2]	TableScan	TAB_DB.PUBLIC.EMPLO	Y]	1	1	3072

EXPLAIN

```
SELECT department, AVG(salary)as department_avgsal
```

```
FROM clone_TAB_DB.PUBLIC.clone_EMP
```

```
GROUP BY department order by department_avgsal desc;
```

Results Results (just now)

Table Chart

Q 5 rows 343ms ↴

#	tep	# id	A parentOperator	A operation	A objects	A alias	A expressions	# partitionsTotal	# partitionsAssigned	# by
1	null	null	null	GlobalStats	null	null	null	1	1	
2	1	0	null	Result	null	null	CLONE_EMP.DEPARTM	null	null	
3	1	1	[0]	Sort	null	null	SCALED_ROUND_INT_D	null	null	
4	1	2	[1]	Aggregate	null	null	aggExprs:[SUM(CLON	null	null	
5	1	3	[2]	TableScan	CLONE_TAB_DB.PUBLI	null	DEPARTMENT, SALARY	1	1	

SELECT

```
query_text,
execution_time,
bytes_scanned

FROM snowflake.account_usage.query_history

WHERE query_text ILIKE '%AVG(salary)%'

ORDER BY start_time DESC;
```

9.3)explain performance similarities/differences

Clone tables:

Same micro-partitions share pannum

Same metadata structure

Snowflake optimizer:

Same execution plan generate pannum

Assignment 10: Cleanup & Retention

10.1) Drop original objects

```
DROP TABLE TAB_DB.PUBLIC.EMPLOYEE;
```

Results ... Results (just now)

Table Chart

Q 1 row 87ms ↴

#	A status
1	EMPLOYEE successfully dropped.

```
select*from tab_db.public.employee; ----doesnt exist
```



10.2) Validate cloned object availability

```
SELECT * FROM CLONE_tab_db.PUBLIC.EMPLOYEE;
```

-- i droped original table but showing clone table no affect

The screenshot shows a database interface with a results table. At the top, there are tabs for 'Results' and 'Results (just now)', with 'Results (just now)' being active. Below the tabs, it says '11 rows 64ms'. The table displays employee data with various columns: EMP_ID, EMP_NAME, DEPARTMENT, JOB_TITLE, SALARY, LOCATION, and HIRE_DATE. The data includes rows for Michael, Ramesh, Meena, Priya, David, Sneha, and Arjun.

#	EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
1	107	Michael	Finance	Account Director	126000	New York	2015-01-12
2	101	Ramesh	IT	Data Engineer	107000	Bangalore	2019-06-15
3	104	Meena	IT	Senior Data Analyst	113000	Bangalore	2018-11-23
4	106	Priya	HR	HR Business Partner	83000	Hyderabad	2019-03-18
5	105	David	IT	Software Engineer	112000	San Francisco	2021-07-05
6	108	Sneha	IT	Junior Developer	90000	Bangalore	2022-08-09
7	109	Arjun	IT	Data Architect	127000	Bangalore	2017-05-20

10.3) Drop clone and analyze Time Travel behavior

```
DROP TABLE CLONE_tab_db.PUBLIC.EMPLOYEE;
```

```
SELECT * FROM CLONE_tab_db.PUBLIC.EMPLOYEE; --- its showing doesnt exist
```



Recover Clone Using Time Travel (UNDROP)

```
undrop table clone_tab_db.public.employee;
```

The screenshot shows a database interface with a results table. At the top, there are tabs for 'Results' and 'Results (just now)', with 'Results (just now)' being active. Below the tabs, it says '1 row 65ms'. The table displays a single row with status information: 'status' and 'Table EMPLOYEE successfully restored.'

#	status
1	Table EMPLOYEE successfully restored.

--Verify Data After UNDROP:

SELECT * FROM CLONE_tab_db.PUBLIC.EMPLOYEE; ---- its showing output

#	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
	101	Emily Ramesh +7 more	IT Marketing +2 more	Marketing Manager Data Engineer +6 more	83000 127000	Bangalore Chicago +3 more	45.5% 27.3% 1/12/2015 8/9/2022
1	107	Michael	Finance	Account Director	126000	New York	2015-01-12
2	101	Ramesh	IT	Data Engineer	107000	Bangalore	2019-06-15
3	104	Meena	IT	Senior Data Analyst	113000	Bangalore	2018-11-23
4	106	Priya	HR	HR Business Partner	83000	Hyderabad	2019-03-18
5	105	David	IT	Software Engineer	112000	San Francisco	2021-07-05
6	108	Sneha	IT	Junior Developer	90000	Bangalore	2022-08-09

10.4) Explain lifecycle and retention rules

Dropped cloned objects can be recovered within the Time Travel period, after which data moves to Fail-safe and is eventually permanently deleted.