

Original table:

..The main source table where data is first stored.

Cloned table:

.. A copy created instantly using zero-copy cloning

Assignment 1: Base Table Analysis

1. Load both datasets into Snowflake tables

Employee table:

```
1      create database emp_table;
2      create schema emp_schema;
3
4      create or replace table employees(
5          emp_id int,
6          emp_name varchar(100),
7          department varchar(100),
8          job_title varchar(100),
9          salary int,
10         location varchar(100),
11         hire_date date
12     );
13
14     INSERT INTO employees VALUES
15     (101,'Ramesh','IT','Data Engineer',95000,'Bangalore','2019-06-15'),
16     (102,'Anita','HR','HR Executive',65000,'Hyderabad','2020-02-10'),
17     (103,'John','Finance','Finance Manager',110000,'New York','2016-09-01'),
18     (104,'Meena','IT','Senior Data Analyst',98000,'Bangalore','2018-11-23'),
19     (105,'David','IT','Software Engineer',90000,'San Francisco','2021-07-05'),
20     (106,'Priya','HR','HR Business Partner',72000,'Hyderabad','2019-03-18'),
21     (107,'Michael','Finance','Account Director',120000,'New York','2015-01-12'),
22     (108,'Sneha','IT','Junior Developer',78000,'Bangalore','2022-08-09'),
23     (109,'Arjun','IT','Data Architect',115000,'Bangalore','2017-05-30'),
24     (110,'Emily','Marketing','Marketing Manager',88000,'Chicago','2018-10-17');
```

Results (just now)	
	Table
000	A status
1 Table EMPLOYEES successfully created.	

Department table:

```
31     CREATE OR REPLACE TABLE departments (
32         dept_id INT,
33         department varchar(100),
34         manager varchar(100),
35         budget int,
36         location varchar(100)
37     );
38
39     INSERT INTO departments (dept_id, department, manager, budget, location) VALUES
40     (10, 'IT', 'Arjun Rao', 500000, 'Bangalore'),
41     (20, 'HR', 'Anita Sharma', 200000, 'Hyderabad'),
42     (30, 'Finance', 'John Smith', 600000, 'New York'),
43     (40, 'Marketing', 'Emily Wilson', 300000, 'Chicago');
```

Results (14 hours ago)	
	Table
000	A status
1 Table DEPARTMENTS successfully created.	

2. Validate row counts and data distribution

Employee table:

```
25
26   SELECT department, COUNT(*) AS emp_count
27   FROM employees
28   GROUP BY department;
29
```

Results (just now)

Table Chart

#	DEPARTMENT	EMP_COUNT
1	Marketing	1
2	IT	5
3	HR	2
4	Finance	2

Department table:

```
44
45   SELECT department, COUNT(*) AS dept_count
46   FROM departments
47   GROUP BY department;
48
```

Results (just now)

Table Chart

#	DEPARTMENT	DEPT_COUNT
1	Marketing	1
2	IT	1
3	HR	1
4	Finance	1

3. Identify candidate tables for cloning

Candidate Tables - source table

1. Employees 2. Departments

4. Explain why these tables are suitable for cloning experiments

- **Small and manageable size**

Both `employees` and `departments` tables have a limited number of rows, making them ideal for quick cloning and easy observation of changes.

- **Stable base data**

The data does not change frequently, which is perfect for testing clone behavior without constant data modifications.

- **Well-structured schema**

Clearly defined columns (IDs, names, departments, locations) help in easily validating cloned data.

- **No complex transformations**

These are base tables (not views or CTAS outputs), fully aligned with the rule of using **only cloning**.

- **Ideal for zero-copy cloning demonstration**

Snowflake clones share underlying storage initially, allowing students to observe how storage is saved until DML operations occur.

Assignment 2: Table-Level Cloning

1. Clone the employee table

```

49
50      CREATE OR REPLACE TABLE employees_clone
51      CLONE employees;
52
Results (just now)
Table Chart
# status
1 Table EMPLOYEES_CLONE successfully created.

```

2. Compare:

- **Row count** : Between clone and original

```

52
53      SELECT COUNT(*) AS clone_count FROM employees_clone;
54
Results (just now)
Table Chart
# CLONE_COUNT
1

```

```

55      SELECT COUNT(*) AS original_count FROM employees;
56
Results (just now)
Table Chart
# ORIGINAL_COUNT
1

```

...Both tables return **10 rows**

...Confirms data consistency between original and clone

- **Query performance:**

57 | SELECT * FROM employees WHERE department = 'IT';
 58

Results (just now)

Table Chart

Q ⚡ 5 rows ⓘ 33ms ⚡ ⌂

#	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	# SALARY	LOCATION	HIRE_DATE
1	101	Ramesh	IT	Data Engineer	95000	Bangalore	2019-06-15
2	104	Meena	IT	Senior Data Analyst	98000	Bangalore	2018-11-23
3	105	David	IT	Software Engineer	90000	San Francisco	2021-07-05
4	108	Sneha	IT	Junior Developer	78000	Bangalore	2022-08-09
5	109	Arjun	IT	Data Architect	115000	Bangalore	2017-05-30

59 | SELECT * FROM employees_clone WHERE department = 'IT';
 60
 ↵

Results (just now)

Table Chart

Q ⚡ 5 rows ⓘ 46ms ⚡ ⌂

#	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	# SALARY	LOCATION	HIRE_DATE
1	101	Ramesh	IT	Data Engineer	95000	Bangalore	2019-06-15
2	104	Meena	IT	Senior Data Analyst	98000	Bangalore	2018-11-23
3	105	David	IT	Software Engineer	90000	San Francisco	2021-07-05
4	108	Sneha	IT	Junior Developer	78000	Bangalore	2022-08-09
5	109	Arjun	IT	Data Architect	115000	Bangalore	2017-05-30

..Query execution time is nearly identical

..Clone uses the same underlying micro-partitions as the original table

- **Storage usage:**

- **BYTES** shows how much physical storage the table is using

61 | SELECT
 62 | table_name,
 63 | bytes,
 64 | row_count
 65 | FROM information_schema.tables
 66 | WHERE table_name IN ('EMPLOYEES', 'EMPLOYEES_CLONE');
 67

Results (just now)

Table Chart

Q ⚡ 2 rows ⓘ 1.6s ⚡ ⌂

#	TABLE_NAME	BYTES	ROW_COUNT
1	EMPLOYEES_CLONE	3072	10
2	EMPLOYEES	3072	10

3. Explain why clone creation is instantaneous

..Snowflake uses **zero-copy cloning**

..No physical data is copied during clone creation

..Only metadata pointers are created

Assignment 3: Change Isolation Testing

....Checking that changes made in one table do NOT affect another table

1. Update salary data in the original employee table

```
69 UPDATE employees
70 SET salary = salary + 5000
71 WHERE department = 'IT';
72
```

Results (just now)

Table		Chart	
00	# number of rows updated	## number of multi-joined rows updated	
1		5	0

```
69 UPDATE employees
70 SET salary = salary + 5000
71 WHERE department = 'IT';
72
73 select*from employees;
```

Results (just now)

Table		Chart	
00	# EMP_ID	EMP_NAME	DEPARTMENT
101	110	Anita Arjun +8 more	IT Finance +2 more
3	106	Priya	HR
4	107	Michael	Finance
5	101	Ramesh	IT
6	104	Meena	IT
7	105	David	IT
8	108	Sneha	IT
9	109	Arjun	IT
10	110	Emily	Marketing

2. Delete records in the cloned table

```
72
73 DELETE FROM employees_clone
74 WHERE department = 'HR';
75
```

Results (just now)

Table		Chart	
00	# number of rows deleted	##	
1		2	

```
73
74 select*from employees_clone;
75
76 DELETE FROM employees_clone
77 WHERE department = 'HR';
```

Results (just now)

Table		Chart	
00	# EMP_ID	EMP_NAME	DEPARTMENT
1	103	John	Finance
2	107	Michael	Finance
3	101	Ramesh	IT
4	104	Meena	IT
5	105	David	IT
6	108	Sneha	IT
7	109	Arjun	IT
8	110	Emily	Marketing
9	111	Rahul	IT

3. Insert new records in the clone

```

76   INSERT INTO employees_clone
77     (emp_id, emp_name, department, job_title, salary, location, hire_date)
78   VALUES
79     (111, 'Rahul', 'IT', 'Cloud Engineer', 105000, 'Bangalore', '2023-01-15');
80

```

Results (just now)	
Table	Chart
00	# number of rows inserted
1	1

```

78   INSERT INTO employees_clone
79     (emp_id, emp_name, department, job_title, salary, location, hire_date)
80   VALUES
81     (111, 'Rahul', 'IT', 'Cloud Engineer', 105000, 'Bangalore', '2023-01-15');
82
83   select*from employees_clone;
84

```

Results (just now)																													
Table	Chart																												
00	<table border="1"> <thead> <tr> <th># EMP_ID</th> <th>EMP_NAME</th> <th>DEPARTMENT</th> <th>JOB_TITLE</th> <th>SALARY</th> <th>LOCATION</th> <th>HIRE_DATE</th> </tr> </thead> <tbody> <tr> <td>101</td> <td>Rahul</td> <td>IT</td> <td>Cloud Engineer</td> <td>83000</td> <td>Bangalore</td> <td>1/12/2015</td> </tr> <tr> <td></td> <td>Arjun</td> <td>Finance</td> <td>Account Director</td> <td>120000</td> <td>New York</td> <td>1/15/2023</td> </tr> <tr> <td></td> <td>+7 more</td> <td>+1 more</td> <td>+7 more</td> <td>+2 more</td> <td></td> <td></td> </tr> </tbody> </table>	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE	101	Rahul	IT	Cloud Engineer	83000	Bangalore	1/12/2015		Arjun	Finance	Account Director	120000	New York	1/15/2023		+7 more	+1 more	+7 more	+2 more		
# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE																							
101	Rahul	IT	Cloud Engineer	83000	Bangalore	1/12/2015																							
	Arjun	Finance	Account Director	120000	New York	1/15/2023																							
	+7 more	+1 more	+7 more	+2 more																									
3	101 Ramesh	IT	Data Engineer	100000	Bangalore	2019-06-15																							
4	104 Meena	IT	Senior Data Analyst	103000	Bangalore	2018-11-23																							
5	105 David	IT	Software Engineer	95000	San Francisco	2021-07-05																							
6	108 Sneha	IT	Junior Developer	83000	Bangalore	2022-08-09																							
7	109 Arjun	IT	Data Architect	120000	Bangalore	2017-05-30																							
8	110 Emily	Marketing	Marketing Manager	88000	Chicago	2018-10-17																							
9	111 Rahul	IT	Cloud Engineer	105000	Bangalore	2023-01-15																							
10	111 Rahul	IT	Cloud Engineer	105000	Bangalore	2023-01-15																							

4. Validate isolation between original and clone

Original table:

```

84
85   SELECT emp_id, salary
86   FROM employees
87   WHERE department = 'IT';
88

```

Results (just now)													
Table	Chart												
00	<table border="1"> <thead> <tr> <th># EMP_ID</th> <th>SALARY</th> </tr> </thead> <tbody> <tr> <td>101</td> <td>110000</td> </tr> <tr> <td>104</td> <td>113000</td> </tr> <tr> <td>105</td> <td>105000</td> </tr> <tr> <td>108</td> <td>93000</td> </tr> <tr> <td>109</td> <td>130000</td> </tr> </tbody> </table>	# EMP_ID	SALARY	101	110000	104	113000	105	105000	108	93000	109	130000
# EMP_ID	SALARY												
101	110000												
104	113000												
105	105000												
108	93000												
109	130000												

```

86
87   SELECT *
88   FROM employees
89   WHERE department = 'HR';
90

```

Results (just now)																						
Table	Chart																					
00	<table border="1"> <thead> <tr> <th># EMP_ID</th> <th>EMP_NAME</th> <th>DEPARTMENT</th> <th>JOB_TITLE</th> <th>SALARY</th> <th>LOCATION</th> <th>HIRE_DATE</th> </tr> </thead> <tbody> <tr> <td>102</td> <td>Anita</td> <td>HR</td> <td>HR Executive</td> <td>65000</td> <td>Hyderabad</td> <td>2020-02-10</td> </tr> <tr> <td>106</td> <td>Priya</td> <td>HR</td> <td>HR Business Partner</td> <td>72000</td> <td>Hyderabad</td> <td>2019-03-18</td> </tr> </tbody> </table>	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE	102	Anita	HR	HR Executive	65000	Hyderabad	2020-02-10	106	Priya	HR	HR Business Partner	72000	Hyderabad	2019-03-18
# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE																
102	Anita	HR	HR Executive	65000	Hyderabad	2020-02-10																
106	Priya	HR	HR Business Partner	72000	Hyderabad	2019-03-18																

Clone table:

```
--  
90 | SELECT emp_id, salary  
91 | FROM employees_clone  
92 | WHERE department = 'IT';  
93 |
```

Results (just now)

Table Chart

6 rows 73ms

#	EMP_ID	SALARY
1	101	100000
2	104	103000
3	105	95000
4	108	83000
5	109	120000
6	111	105000

```
/2  
73 | select*from employees_clone;  
74 |  
75 | DELETE FROM employees_clone  
76 | WHERE department = 'HR';  
77 |
```

Results (just now)

Table Chart

10 rows 65ms

#	EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
1	103	John	Finance	Finance Manager	110000	New York	2016-09-01
2	107	Michael	Finance	Account Director	120000	New York	2015-01-12
3	101	Ramesh	IT	Data Engineer	100000	Bangalore	2019-06-15
4	104	Meena	IT	Senior Data Analyst	103000	Bangalore	2018-11-23
5	105	David	IT	Software Engineer	95000	San Francisco	2021-07-05
6	108	Sneha	IT	Junior Developer	83000	Bangalore	2022-08-09
7	109	Arjun	IT	Data Architect	120000	Bangalore	2017-05-30
8	110	Emily	Marketing	Marketing Manager	88000	Chicago	2018-10-17

- Update data in the **original table** → clone data remains **unchanged**
- Insert or delete data in the **clone** → original table remains **unchanged**
- Both tables show **different results after changes**

5. Explain backend copy-on-write behavior

Copy-on-write means Snowflake copies data only when changes happen, not when the clone is created

1. Original table and clone **initially share the same data**
2. **No data is copied** at clone creation
3. When a row is **updated, deleted, or inserted**:
Snowflake **creates a new copy only for the changed data**
4. **Unchanged data remains shared**
5. This ensures **data isolation with minimal storage usage**

Assignment 4: Storage Growth Exploration

Storage Growth Exploration means checking when and why storage increases while working with cloned tables.

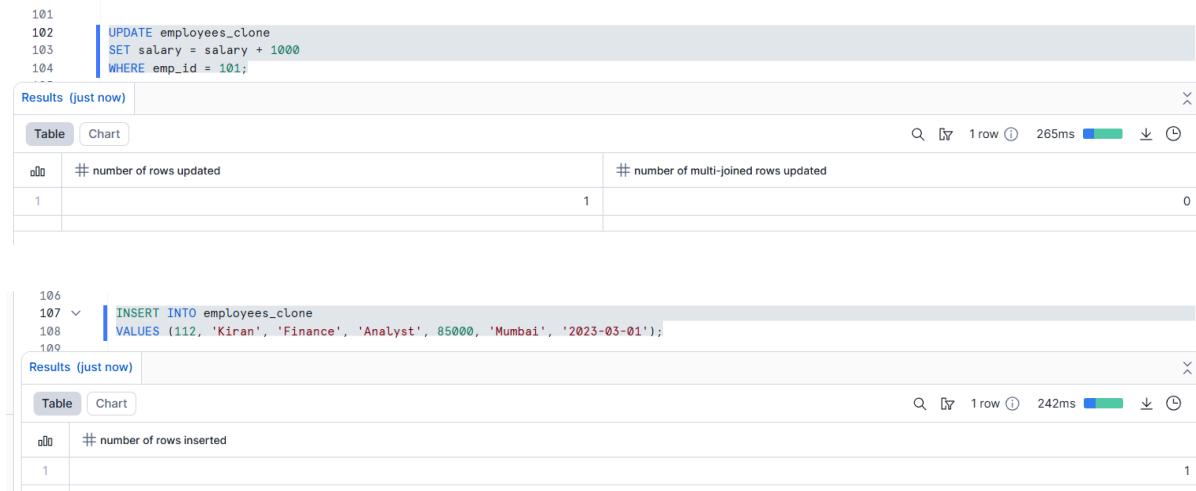
1. Perform DML operations incrementally

Meaning:

Do data changes **one by one**, not all at once.

Example:

- First do **UPDATE**
- Then check storage
- Then do **INSERT**
- Then check storage again



Final Output:

The figure shows the final output of a SELECT query (lines 97-98) on the employees_clone table:

```
97
98 SELECT * FROM employees_clone
```

The results table displays 11 rows of employee data, including columns for EMP_ID, EMP_NAME, DEPARTMENT, JOB_TITLE, SALARY, LOCATION, and HIRE_DATE. The data includes various roles like Cloud Engineer, Account Director, Data Engineer, etc., across departments like IT, Finance, and Marketing, located in cities like Bangalore, New York, and Mumbai, hired between 2012 and 2023.

	# EMP_ID	EMP_NAME	DEPARTMENT	JOB_TITLE	SALARY	LOCATION	HIRE_DATE
1	101	Rahul	IT	Cloud Engineer	83000	Bangalore	2012-01-15
2	112	Arjun	Finance	Account Director	120000	New York	2013-01-15
3	101	Ramesh	IT	Data Engineer	102000	Bangalore	2019-06-15
4	104	Meena	IT	Senior Data Analyst	103000	Bangalore	2018-11-23
5	105	David	IT	Software Engineer	95000	San Francisco	2021-07-05
6	108	Sneha	IT	Junior Developer	83000	Bangalore	2022-08-09
7	109	Arjun	IT	Data Architect	120000	Bangalore	2017-05-30
8	110	Emily	Marketing	Marketing Manager	88000	Chicago	2018-10-17
9	111	Rahul	IT	Cloud Engineer	105000	Bangalore	2023-01-15
10	111	Rahul	IT	Cloud Engineer	105000	Bangalore	2023-01-15
11	112	Kiran	Finance	Analyst	85000	Mumbai	2023-03-01

2. Track storage before and after each operation

```
110 ~ | SELECT table_name, bytes
111 | FROM information_schema.tables
112 | WHERE table_name IN ('EMPLOYEES', 'EMPLOYEES_CLONE');
113 |
Results (just now) X
Table Chart 🔍 ⏪ 2 rows ⓘ 970ms ⚡ ⌂ ⌂

| TABLE_NAME      | BYTES |
|-----------------|-------|
| EMPLOYEES       | 3072  |
| EMPLOYEES_CLONE | 3072  |


```

Before any DML (UPDATE / INSERT / DELETE):

Check storage → **no increase**

After UPDATE:

Check storage again → **storage increases slightly**

After INSERT:

Check storage again → **storage increases more**

3. Identify the exact moment storage increases

Storage increases at the moment when a DML operation is executed (UPDATE, INSERT, or DELETE).

- Creating a clone → **no storage increase**
- Running SELECT queries → **no storage increase**
- First UPDATE / INSERT / DELETE → **storage increases**

4. Explain Snowflake's physical storage strategy

- Snowflake stores data in **small read-only blocks** called **micro-partitions**
- These micro-partitions are **compressed and optimized** automatically
- When a table is cloned, **data is shared**, not copied
- Data blocks are **never changed directly**
- When data is updated, inserted, or deleted:
 - ..Snowflake **creates new micro-partitions** for the changed data
 - ..Old micro-partitions remain unchanged
- This method is called **copy-on-write**

Assignment 5: Time Travel Cloning

1. Perform multiple updates on the original table

```
114 UPDATE employees
115 SET salary = salary + 3000
116 WHERE department = 'IT';
117
Results (just now)
Table Chart
# number of rows updated # number of multi-joined rows updated
1 5 0
```

```
117
118 UPDATE employees
119 SET salary = salary + 2000
120 WHERE department = 'HR';
121
Results (just now)
Table Chart
# number of rows updated # number of multi-joined rows updated
1 2 0
```

Final Output:

```
114 UPDATE employees
115 SET salary = salary + 3000
116 WHERE department = 'IT';
117
118 UPDATE employees
119 SET salary = salary + 2000
120 WHERE department = 'HR';
121
122 SELECT * FROM employees;
123
Results (just now)
Table Chart
# EMP_ID EMP_NAME DEPARTMENT JOB_TITLE SALARY LOCATION HIRE_DATE
101 Anita IT Account Director 67000 Bangalore 2016-09-01
102 Arjun Finance Data Architect 113000 Hyderabad 2020-02-10
103 John Finance Manager 67000 New York 2019-03-18
104 Priya HR Executive 74000 Hyderabad 2015-01-12
105 Michael HR Business Partner 120000 New York 2019-06-15
106 Ramesh Finance Account Director 116000 Bangalore 2018-11-23
107 Meena IT Senior Data Analyst 108000 San Francisco 2021-07-05
108 David IT Software Engineer 96000 Bangalore 2022-08-09
109 Sneha IT Junior Developer 113000 Bangalore 8/9/2022
```

2. Create clones from:

- Current state

```
124 CREATE OR REPLACE TABLE employees_clone_current
125 CLONE employees;
126
127
Results (just now)
Table Chart
status
1 Table EMPLOYEES_CLONE_CURRENT successfully created.
```

- Past state (Time Travel)

```

127 CREATE OR REPLACE TABLE employees_clone_past
128   CLONE employees
129   AT (OFFSET => -300);
130
131
132
133
134
135
136
137

```

results (just now) | X

Table Chart

1 Table EMPLOYEES_CLONE_PAST successfully created.

Q D 1 row 691ms ↴ ⌂

3. Compare data across versions

```

131   SELECT
132     c.emp_id,
133     c.salary AS current_salary,
134     p.salary AS past_salary
135   FROM employees_clone_current c
136   JOIN employees_clone_past p
137   ON c.emp_id = p.emp_id;
138

```

results (2 minutes ago) | X

Table Chart

10 rows 172ms ↴ ⌂

#	EMP_ID	CURRENT_SALARY	PAST_SALARY
1	103	110000	110000
2	102	67000	67000
3	106	74000	74000
4	107	120000	120000
5	101	113000	113000
6	104	116000	116000
7	105	108000	108000
8	108	96000	96000
9	109	133000	133000

4. Explain real-world recovery use cases

- Recover data after **accidental UPDATE or DELETE**
- Restore a table to a **previous correct version**
- Compare data **before and after a mistake**
- Quickly fix **production issues** without reloading data
- Support **audit and investigation** needs

Assignment 6: Schema-Level Cloning

1. Clone a schema containing employee and department tables

```
144 CREATE OR REPLACE SCHEMA hr_schema_clone
145 CLONE emp_schema;
146
147
Results (just now)  
Table Chart   1 row  2.1s   

| id | status                                       |
|----|----------------------------------------------|
| 1  | Schema HR_SCHEMA_CLONE successfully created. |


```

2. Modify tables in cloned schema

Employee table:

```
148 UPDATE hr_schema_clone.employees
149 SET salary = salary + 3000
150 WHERE department = 'IT';
151
Results (just now)  
Table Chart   1 row  257ms   

| id | # number of rows updated | # number of multi-joined rows updated |
|----|--------------------------|---------------------------------------|
| 1  | 5                        | 0                                     |


```

```
155 select * from hr_schema_clone.employees;
156
Results (just now)  
Table Chart   10 rows  106ms   

| id | # EMP_ID | EMP_NAME | DEPARTMENT | JOB_TITLE           | SALARY | LOCATION      | HIRE_DATE  |
|----|----------|----------|------------|---------------------|--------|---------------|------------|
| 2  | 102      | Anita    | HR         | HR Executive        | 67000  | Hyderabad     | 2020-02-10 |
| 3  | 106      | Priya    | HR         | HR Business Partner | 74000  | Hyderabad     | 2019-03-18 |
| 4  | 107      | Michael  | Finance    | Account Director    | 120000 | New York      | 2015-01-12 |
| 5  | 101      | Ramesh   | IT         | Data Engineer       | 119000 | Bangalore     | 2019-06-15 |
| 6  | 104      | Meena    | IT         | Senior Data Analyst | 122000 | Bangalore     | 2018-11-23 |
| 7  | 105      | David    | IT         | Software Engineer   | 114000 | San Francisco | 2021-07-05 |
| 8  | 108      | Sneha    | IT         | Junior Developer    | 102000 | Bangalore     | 2022-08-09 |
| 9  | 109      | Arjun    | IT         | Data Architect      | 139000 | Bangalore     | 2017-05-30 |
| 10 | 110      | Emily    | Marketing  | Marketing Manager   | 88000  | Chicago       | 2018-10-17 |


```

Department table:

```
152 INSERT INTO hr_schema_clone.departments
153 VALUES (50, 'Sales', 'Rahul Kumar', 250000, 'Mumbai');
154
Results (just now)  
Table Chart   1 row  240ms   

| id | # number of rows inserted |
|----|---------------------------|
| 1  | 1                         |


```

155 select * from hr_schema_clone.departments;

156

Results (just now)

Table Chart

#	DEPT_ID	DEPARTMENT	MANAGER	BUDGET	LOCATION
1	10	IT	Arjun Rao	500000	Bangalore
2	20	HR	Anita Sharma	200000	Hyderabad
3	30	Finance	John Smith	600000	New York
4	40	Marketing	Emily Wilson	300000	Chicago
5	50	Sales	Rahul Kumar	250000	Mumbai
6	50	Sales	Rahul Kumar	250000	Mumbai

3. Drop tables in original schema

Departments table:

156

157 DROP TABLE hr_schema_clone.departments;

158

Results (just now)

Table Chart

#	status
1	DEPARTMENTS successfully dropped.

Employees table:

159

160 DROP TABLE hr_schema_clone.employees;

161

Results (just now)

Table Chart

#	status
1	EMPLOYEES successfully dropped.

4. Observe behavior and explain metadata inheritance

Metadata inheritance means the clone gets its own metadata while sharing data blocks until changes occur.

Observed behavior:

- Changes made in the cloned schema do not affect the original schema
- Dropping tables in the original schema does not delete tables in the clone
- Both schemas work independently

Metadata inheritance (meaning):

- When a schema is cloned, Snowflake copies metadata only, not the actual data
- Both schemas initially point to the same underlying data blocks
- Each schema has its own metadata
- Any change creates new data using copy-on-write

Assignment 7: Database-Level Cloning

1. Clone the entire database

```
161 | CREATE OR REPLACE DATABASE emp_db_clone  
162 |   CLONE my_db;  
163
```

Results (just now)	
	Table
000	A status
1 Database FMP DR CLONE successfully created.	

2. Test query access and object availability

```
165 | USE DATABASE emp_db_clone;
```

Results (just now)	
	Table
000	A status
1 Statement executed successfully.	

```
166 | SHOW SCHEMAS;
```

Results (just now)							
	Table	Chart	Q	Y	3 rows	1	40ms
000	created_on	A name	A is_default	A is_current	A database_name	A owner	A comment
1	2025-12-22 07:41:51.176 -0800	INFORMATION_SCHEMA	N	N	EMP_DB_CLONE		Views describing the contents of schemas in the database.
2	2025-12-22 07:35:48.197 -0800	MY_SCHEMA	N	N	EMP_DB_CLONE	ACCOUNTADMIN	
3	2025-12-22 07:35:48.197 -0800	PUBLIC	N	Y	EMP_DB_CLONE	ACCOUNTADMIN	

```
167 | SELECT * FROM hr_schema_clone.employees;
```

Results (just now)							
	Table	Chart	Q	Y	10 rows	1	765ms
000	# EMP_ID	A EMP_NAME	A DEPARTMENT	A JOB_TITLE	# SALARY	A LOCATION	HIRE_DATE
1	103	John	Finance	Finance Manager	110000	New York	2016-09-01
2	102	Anita	HR	HR Executive	67000	Hyderabad	2020-02-10
3	106	Priya	HR	HR Business Partner	74000	Hyderabad	2019-03-18
4	107	Michael	Finance	Account Director	120000	New York	2015-01-12
5	101	Ramesh	IT	Data Engineer	113000	Bangalore	2019-06-15
6	104	Meena	IT	Senior Data Analyst	116000	Bangalore	2018-11-23
7	105	David	IT	Software Engineer	108000	San Francisco	2021-07-05
8	108	Sneha	IT	Junior Developer	96000	Bangalore	2022-08-09
9	109	Arjun	IT	Data Architect	133000	Bangalore	2017-05-30

- All schemas and tables are available
- Data can be queried normally

3. Perform DML on cloned database

The screenshot shows two separate DML operations in the Snowflake interface:

```
169 UPDATE hr_schema_clone.employees
170 SET salary = salary + 4000
171 WHERE department = 'IT';
172
```

Results (just now)

Table	Chart
# number of rows updated	# number of multi-joined rows updated
1	5
	0


```
173 INSERT INTO hr_schema_clone.departments
174 VALUES (60, 'Support', 'Neha Singh', 180000, 'Pune');
175
```

Results (just now)

Table	Chart
# number of rows inserted	
1	1

- Changes apply **only to the cloned database**
- Original database remains unchanged

4. Explain how Snowflake isolates environments (Dev/Test/Prod)

- Snowflake uses **zero-copy database cloning**
- Dev, Test, and Prod databases start by **sharing the same data**
- Each environment has **independent metadata**
- Changes in Dev/Test do **not affect Prod**
- New data is created only when changes happen (**copy-on-write**)

Assignment 9: Performance Comparison

1. Run identical analytical queries on:
 - Original tables

The screenshot shows an analytical query running on original tables:

```
182
183     department,
184         AVG(salary) AS avg_salary
185 FROM employees
186 GROUP BY department;
```

Res (13 hours ago) | Results (just now)

Table	Chart
# DEPARTMENT	# AVG_SALARY
1 IT	95200.000000
2 HR	68500.000000
3 Finance	115000.000000
4 Marketing	88000.000000

- Cloned tables

```

188
189     department,
190     AVG(salary) AS avg_salary
191
192 FROM employees_clone
193 GROUP BY department;
194

```

Res (12 hours ago) Results (just now)

Table Chart

#	DEPARTMENT	AVG_SALARY
1	IT	95200.000000
2	HR	68500.000000
3	Finance	115000.000000
4	Marketing	88000.000000

Both queries return the same result.

2. Compare execution plans

Explain using text - we can identify the query performance.

Original table:

```

194 EXPLAIN USING TEXT
195
196 SELECT
197     department,
198     AVG(salary)
199 FROM employees
200 GROUP BY department;
201

```

Res (13 hours ago) Results (just now)

Table Chart

#	content
1	GlobalStats: partitionsTotal=1 partitionsAssigned=1 bytesAssigned=3072 Operations: 1:0 →Result EMPLOYEES.DEPARTMENT, SCALED_ROUND_INT_DIVIDE(SUM(EMPLOYEES.SALARY), COU

Clone table:

```

200
201 EXPLAIN USING TEXT
202 SELECT
203     department,
204     AVG(salary)
205 FROM employees_clone
206 GROUP BY department;
207

```

Res (13 hours ago) Results (just now)

Table Chart

#	content
1	GlobalStats: partitionsTotal=1 partitionsAssigned=1 bytesAssigned=3072 Operations: 1:0 →Result EMPLOYEES_CLONE.DEPARTMENT, SCALED_ROUND_INT_DIVIDE(SUM(EMPLOYEES_CLONE

- Execution plans are **nearly identical**
- Same micro-partitions are scanned
- Same query optimization is applied

3. Explain performance similarities/differences

Performance is almost the same because original and cloned tables share the same data and execution plan; differences appear only after data changes or cache effects.

Assignment 10: Cleanup & Retention

1. Drop original objects

Employees:

```
208 | DROP TABLE employees;
209 | Results (just now)
Table | Chart | X
1 EMPLOYEES successfully dropped.
```

Departments:

```
210 | DROP TABLE departments;
211 | Results (just now)
Table | Chart | X
1 DEPARTMENTS successfully dropped.
```

2. Validate cloned object availability

```
212 | SELECT * FROM employees_clone;
213 | Results (just now)
Table | Chart | X
1 EMP_ID EMP_NAME DEPARTMENT JOB_TITLE SALARY LOCATION HIRE_DATE
1 101 Ramesh IT Data Engineer 95000 Bangalore 2019-06-15
2 102 Anita HR HR Executive 65000 Hyderabad 2020-02-10
3 103 John Finance Finance Manager 110000 New York 2016-09-01
4 104 Meena IT Senior Data Analyst 98000 Bangalore 2018-11-23
5 105 David IT Software Engineer 90000 San Francisco 2021-07-05
6 106 Priya HR HR Business Partner 72000 Hyderabad 2019-03-18
7 107 Michael Finance Account Director 120000 New York 2015-01-12
```

```
214 | SELECT * FROM departments_clone;
215 | Results (just now)
Table | Chart | X
1 DEPT_ID DEPARTMENT MANAGER BUDGET LOCATION
1 10 IT Arjun Rao 500000 Bangalore
2 20 HR Anita Sharma 200000 Hyderabad
3 30 Finance John Smith 600000 New York
4 40 Marketing Emily Wilson 300000 Chicago
```

3. Drop clone and analyze Time Travel behavior

Drop employees_clone / By using time travel the data is available. After retention period ends, data is permanently removed:

```
215 | DROP TABLE employees_clone;
216 |
217 | Res (13 hours ago) Results (just now)
| Table Chart
| 00  A status
| 1 EMPLOYEES_CLONE successfully dropped.
```

```
217 | SELECT * FROM employees_clone
218 | AT (OFFSET => -60);
219 |
Res (13 hours ago) Results (just now)
Table Chart
00  # EMP_ID  A EMP_NAME  ;  A DEPARTMENT  A JOB_TITLE  # SALARY  A LOCATION  HIRE_DATE
1    101   Ramesh    IT      Data Engineer    95000  Bangalore  2019-06-15
2    102   Anita     HR      HR Executive    65000  Hyderabad  2020-02-10
3    103   John      Finance  Finance Manager  110000 New York  2016-09-01
4    104   Meena    IT      Senior Data Analyst 98000  Bangalore  2018-11-23
5    105   David    IT      Software Engineer  90000  San Francisco  2021-07-05
6    106   Priya    HR      HR Business Partner 72000  Hyderabad  2019-03-18
7    107   Michael  Finance Account Director  120000 New York  2015-01-12
```

4. Explain lifecycle and retention rules

- Dropped objects go into **Time Travel**
- Data can be recovered within **retention period**
- After Time Travel → data moves to **Fail-safe**
- After Fail-safe → data is **permanently deleted**
- Retention period depends on **Snowflake edition**